

Frontier Max

From benchmark to runtime.

Turn public AI benchmark evidence into the right model policy for the work.

agent-frontier.monilpat.chatgpt.site

EVIDENCE



POLICY



ROUTE



Leaderboards answer the wrong question.

“Which model scored highest?”

A universal rank compresses away the conditions that determine whether a model actually works in production.

“Which configuration can meet this work’s quality, time, cost, and context limits?”



QUALITY TIME COST CONTEXT

One rank hides four constraints.

“Best” changes when the workload changes.



TIME

Human waiting
Interactive coding



COST

Agent delegated
Long-horizon work

OTHER OBJECTIVES

REAL-TIME

p95 latency

BATCH

accepted
throughput

CONTINUOUS

reliability + ops

**An agent is a curve,
not a point.**

One transparent chain turns evidence into action.

READ

01

BenchmarkList catalog
Source + version lineage

DECIDE

02

Workload-specific
frontier
Hard gates before trade-
offs

ROUTE

03

OpenRouter ×
OpenCode
Reproducible session
policy

PROVE

04

Content-free receipt
Policy + evidence
hashes

No universal weighted rank. No hidden proxy. No pay-to-win placement.

Describe the work—not the model.

INTERACTIVE

“Fix a flaky authentication test while I pair.”

code.interactive

Current route proxy:
fastest p50 throughput in the high coding tier

DELEGATED


“Migrate thirty endpoints overnight and open a verified PR.”

code.delegated

Current route proxy:
cheapest available model in the high coding tier

Same evidence universe. Different objective. Different route.

Gemini understands intent. Inspectable code decides.



Structured output · explicit assumptions · content-free receipt

Only the short routing brief reaches Gemini. Repository contents, code, diffs, keys, and downstream conversations do not.

Free evidence earns trust. Decisioning captures value.

WEDGE

Coding-agent model selection

A visible difference between
interactive and delegated work.

BUYER

AI platform, engineering, FinOps

Teams spending real money across
multiple models and providers.

REVENUE

Policies, alerts, governance, routing API

Shared decisions and audit-ready selection
inside production workflows.

DURABLE MOAT

Historical provenance + the workload policies teams operationalize

The benchmark layer is becoming a control plane.

Help turn every benchmark into a decision teams can defend—and execute.

BUILT BY

Monil Patel · Staff Applied AI · 7 years at Meta

Evaluation, post-training, and coding-agent systems

WE'RE LOOKING FOR

05 design partners

Costly multi-model workflows

GCP product support

Gemini decision layer

VC investor partners

A durable model-selection control plane

Frontier Max

agent-frontier.monilpat.chatgpt.site