

From benchmark to runtime.

Frontier Max turns benchmark data into the right model for the work—explaining the evidence, applying workload-aware trade-offs, and launching a reproducible route.

THE GAP A leaderboard can tell you who scored highest. It cannot tell you which model meets this task's quality floor, response-time need, context budget, and dollar constraint—or turn that decision into a command.

THE PROBLEM

Benchmarks are plentiful, but legibility and action are missing. Teams stitch together charts, trust a universal #1, then choose a model by hand—collapsing a conditional decision into one score.

THE INSIGHT

"Best" changes with the workload. In interactive coding, latency and throughput become binding after quality gates. For long-horizon agents, cost dominates once latency and token capacity are reasonable. Tokens are always measured, but optimized only when they bind.

Frontier Max treats selection as a workload-conditioned 4D Pareto problem: cost × latency × token capacity × task score.

FOUR LAYERS

EVIDENCE Ingest public model and benchmark sources, including BenchmarkList, while preserving source, version, and provenance.

INTERPRETATION Convert raw metrics into conditional Pareto frontiers—without a universal weighted ranking.

ACTION Produce an executable policy and launch through the open CLI workflow using OpenRouter × OpenCode.

ACCOUNTABILITY Emit content-free local receipts for the selected policy, requested route, and source/version provenance.

WHY IT IS DIFFERENT

Comparison products end with a chart. Frontier Max ends with a policy, a route, and a receipt.

THE TECHNICAL WEDGE

Gemini 3.5 Flash maps natural-language task intent to one of two allowed workload profiles and returns an explanation, signals, confidence, and a caveat. Deterministic validation resolves the checked-in OpenRouter route, including the configured high coding-tier threshold. The model interprets intent; code owns the route.

CURRENT CLI TRUTH Within the high Artificial Analysis coding tier, interactive mode requests OpenRouter's fastest available option by p50 throughput; delegated mode requests its cheapest available option by model price. These are transparent proxies—not accepted-result or cost-per-verified-success claims.

A 60-SECOND PROOF

PAIRING "Fix a flaky authentication test while I pair." → interactive policy → latency/throughput route

OVERNIGHT "Migrate thirty endpoints and open a verified PR." → delegated policy → cost-first route after gates

Same benchmark universe. Different policy. Reproducible receipt.

MARKET & MODEL

Buyer AI platform, engineering productivity, and FinOps teams managing multi-model stacks. **Model** Free evidence catalog + open CLI; paid workspace/API for governance, policies, audit history, alerts, and routing analytics. Frontier Max consumes benchmark data; it does not alter scores or sell placement.

WHY NOW

Models, prices, providers, and context limits move weekly. Policy + provenance lets teams update the route without rewriting the decision.

TEAM & ASK

MONIL PATEL 7 years at Meta; Staff Applied AI; evals, post-training, and coding agents.

THE ASK Five design partners, GCP support to harden the Gemini/Cloud Run path, and investor partners for the decision layer between benchmarks and runtime.